

# Data Integration and Large Scale Analysis

## 02 Data Warehousing and ETL

**Shafaq Siddiqi**

Graz University of Technology, Austria



# Announcements/Org

## ■ #1 Video Recording

- Link in **TUbe** & **TeachCenter**
- Optional attendance
- In-person and video-recorded lectures
  - **HS i5** or Webex: <https://tugraz.webex.com/meet/shafaq.siddiqi>

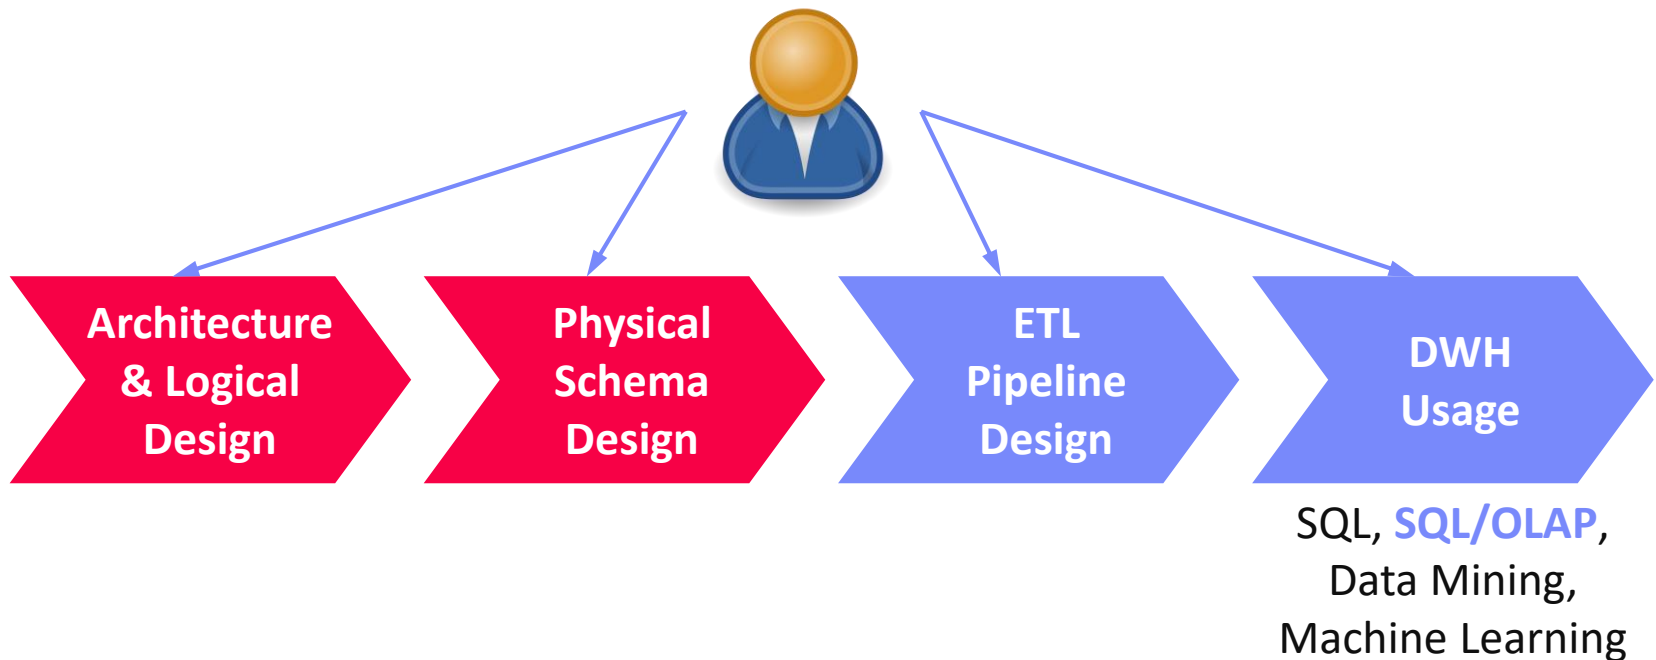


## ■ WKO Research Grants


- <https://www.tugraz.at/en/research/research-at-tu-graz/services-fuer-forschende/foerderprogramme-und-preise-an-der-tu-graz/#c87088>
- Submission deadline October 20, 2023

# Agenda

- **Data Warehousing (DWH)**
- **Extraction, Transformation, Loading (ETL)**
- **SQL/OLAP Extensions**



# Data Warehousing



1. [**Wolfgang Lehner**: Datenbanktechnologie für Data-Warehouse-Systeme. Konzepte und Methoden, Dpunkt Verlag, 1-373, 2003]
2. [**C. S. Jensen, T. B. Pedersen, C. Thomsen**. Multidimensional Databases and Data Warehousing. Morgan and Claypool Publishers. 2010]



# Motivation and Tradeoffs

- **Goal:** Queries over consolidated and cleaned data of several, potentially heterogeneous, data sources



**OLTP** (Online Transaction Processing)  
**vs OLAP** (Online Analytical Processing)




SCM



Material

ERP



CRM

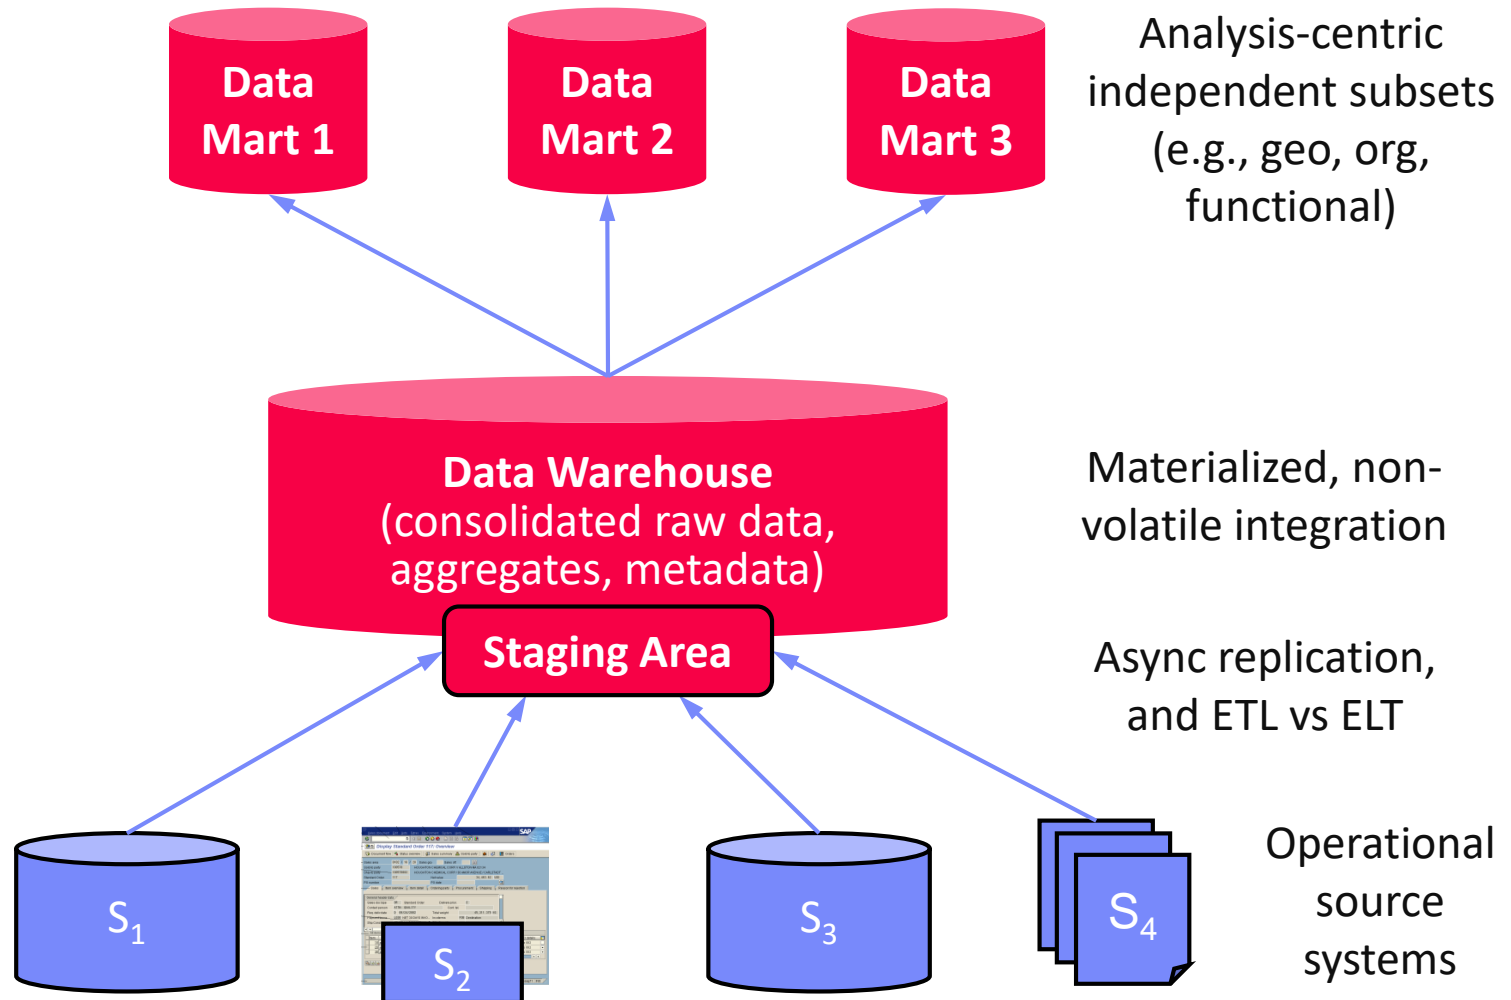


eCommerce

Operational Systems

- **Tradeoffs**
  - **Analytical query performance:** write vs read optimized data stores
  - **Virtualization:** overhead of remote access, source systems affected
  - **Consistency:** sync vs async changes, time regime → up-to-date?
  - **Others:** history, **flexibility**, **redundancy**, effort for **data exchange**


# Data Warehouse Architecture



# Data Warehouse Architecture, cont.

## ■ Data Warehouse (DWH)

- “A data warehouse is a **subject-oriented, integrated, time-varying, non-volatile** collection of data in support of the management's decision-making process.” (Bill Inmon)
- **#1 Subject-oriented:** analysis-centric organization (e.g., sales) → Data Mart
- **#2 Integrated:** consistent data from different data sources
- **#3 Time-varying:** History (snapshots of sources), and temporal modelling
- **#4 Non-volatile:** Read-only access, limited to periodic data loading by admin



## ■ Different DWH Instantiations

- **Single DWH system** with virtual/materialized views for data marts
- Separate systems for consolidated DWH and aggregates/data marts (**dependent data marts**)
- Data-Mart-local staging areas and ETL (**independent data marts**)

# Multi-dimensional Modeling: Data Cube

## ■ Central Metaphor: **Data Cube**



- Qualifying data (categories, dimensions)
- Quantifying data (cells)
- Often sparse (0 for empty cells)

## ■ Multi-dimensional Schema

- Set of **dimension hierarchies** ( $D^1, \dots, D^n$ )
- Set of **measures** ( $M^1, \dots, M^m$ )

## ■ Dimension Hierarchy

- Partially-ordered set  $D$  of categorical attributes ( $\{D_1, \dots, D_n, Top_D\}; \rightarrow$ )
- Generic **maximum element**  
 $\forall i (1 \leq i \leq n): D_i \rightarrow Top_D$
- Existing **minimum element** (primary attribute)  
 $\exists i (1 \leq i \leq n) \forall j (1 \leq j \leq n, i \neq j): D_i \rightarrow D_j$





# Multi-dimensional Modeling: Data Cube, cont.

## ■ Dimension Hierarchy, cont.

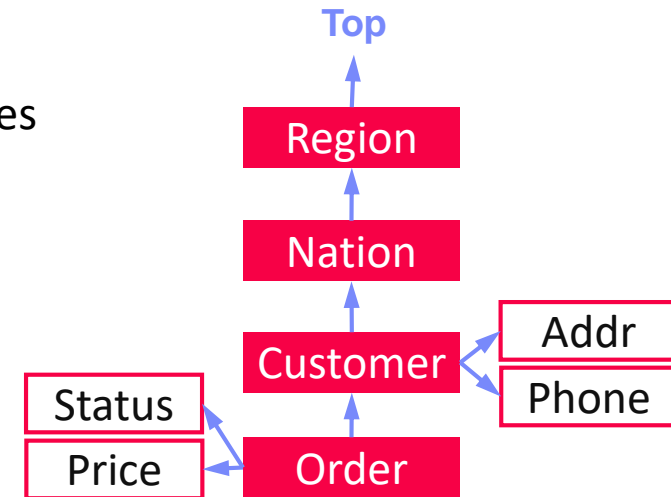
- Classifying (categorical) vs descriptive attributes
- **Orthogonal dimensions:** there are no functional dependencies between attributes of different dimensions

## ■ Fact F

- Base tuples w/ measures of summation type
- Granularity G as subset of categorical attributes

## ■ Measure M


- Computation function over non-empty subset of facts  $f(F_1, \dots, F_k)$  in schema
- Scalar function vs aggregation function
- Granularity G as subset of categorical attributes



# Multi-dimensional Modeling: Operations

## ■ Slicing


- Select a “slice” of the cube by specifying a filter condition on **one of the dimensions** (categorical attributes)
- Same data granularity but subset of dimensions



## ■ Dicing


- Select a “sub-cube” by specifying a filter condition on **multiple dimensions**
- Complex Boolean expressions possible
- Sometimes slicing used synonym

**Example:** Location=Graz **AND**  
Color=White **AND** Make=BMW



# Multi-dimensional Modeling: Operations, cont.

- **Roll-up** (similar Merge - remove dim)
  - Aggregation of facts or measures into coarser-grained aggregates (measures)
  - Same dimensions but different granularity
- **Drill-Down** (similar Split add dim)
  - Disaggregation of measures into finer-grained measures




# Multi-dimensional Modeling: Operations, cont.

## Drill-Across

- Change from one cube to another

## Drill-Through



- Drill-Down to smallest granularity of underlying data store (e.g., RDBMS)
- E.g., find relational tuples



FName	LName	Local	Make	Color
Matthias	Boehm	Graz	BMW	White
...	...	...	...	...

## Pivot

- Rotate cube by exchanging dimensions



# Aggregation Types

## Recap: Classification of Aggregates

- **Additive** aggregation functions (**SUM**, **COUNT**)
- **Semi-additive** aggregation functions (**MIN**, **MAX**)
- **Additively computable** aggregation functions (**AVG**, **STDDEV**, **VAR**)
- Aggregation functions (**MEDIAN**, **QUANTILES**)

## Summation Types of Measures

- **FLOW**: arbitrary aggregation possible
- **STOCK**: aggregation possible, except over temporal dim
- **VPU**: value-per-unit typically (e.g., price)

[Hans-Joachim Lenz, Arie Shoshani:  
Summarizability in OLAP and  
Statistical Data Bases. SSDBM 1997]



[TUGraz online]

## Necessary Conditions

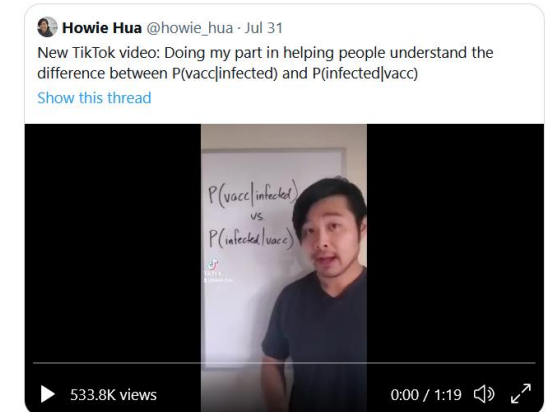

- Disjoint attribute values
- Completeness
- Type compatibility

# Stud	16/17	17/18	18/19	19/20	20/21	Total
CS	1,153	1,283	1,321	1,343	1368	?
SEM	928	970	939	944	985	?
ICE	804	868	846	842	849	?
Total	2,885	3,121	3,106	3,129	3,202	?

# Excursus: Other **Misleading** Statistics

## ■ Problem Setting

- 100 people (**90 vaccinated**, **10 non-vaccinated**)
- 5 infected vaccinated, 2 infected non-vaccinated



[[https://twitter.com/howie\\_hua/status/1421502809862664197](https://twitter.com/howie_hua/status/1421502809862664197)]

- $P(\text{vacc}|\text{infected}) = 5/7 = 0.71 \rightarrow \text{misleading}$
- $P(\text{infected}|\text{vacc}) = 5/90 = 0.056$
- $P(\text{infected}|\text{non-vacc}) = 2/10 = 0.2$

[see also  
Simpson's Paradox  
in **06 Data Cleaning**]

# Aggregation Types, cont.

## ■ Additivity

	FLOW	STOCK: Temporal Agg?		VPU
		Yes	No	
MIN/MAX	✓		✓	✓
SUM	✓	X	✓	X
AVG	✓		✓	✓
COUNT	✓		✓	✓

## ■ Type Compatibility (**addition**/ subtraction)

	FLOW	STOCK	VPU
FLOW	FLOW	STOCK	X
STOCK		STOCK	X
VPU			VPU

# Data Cube Mapping and MDX

## ■ MOLAP (Multi-Dim. OLAP)

- OLAP server with native multi-dimensional data storage
- Dedicated query language: Multidimensional Expressions (MDX)
- E.g., IBM Cognos Powerplay, Essbase

[\[https://docs.microsoft.com/en-us/analysis-services/multidimensional-models/mdx\]](https://docs.microsoft.com/en-us/analysis-services/multidimensional-models/mdx)

```
SELECT
    {[Measures].[Sales],
    [Measures].[Tax]} ON COLUMNS,
    {[Date].[Fiscal].[Year].&[2002],
    [Date].[Fiscal].[Year].&[2003] } ON ROWS
FROM [Adventure Works]
WHERE ([Sales Territory].[Southwest])
```

## ■ ROLAP (Relation OLAP)

- OLAP server w/ storage in RDBMS
- E.g., all commercial RDBMS vendors

## ■ HOLAP (Hybrid OLAP)

- OLAP server w/ storage in RDBMS and multi-dimensional in-memory caches and data structures

**Requires mapping to relational model**

[Example systems:  
[https://en.wikipedia.org/wiki/Comparison\\_of\\_OLAP\\_servers](https://en.wikipedia.org/wiki/Comparison_of_OLAP_servers)]



# Recap: Relational Data Model

- Domain D (value domain): e.g., Set S, INT, Char[20]

- Relation R

- Relation schema RS:  
Set of k attributes  $\{A_1, \dots, A_k\}$
- Attribute  $A_j$ : value domain  $D_j = \text{dom}(A_j)$
- Relation: subset of the Cartesian product over all value domains  $D_j$   
 $R \subseteq D_1 \times D_2 \times \dots \times D_k, k \geq 1$

Attribute

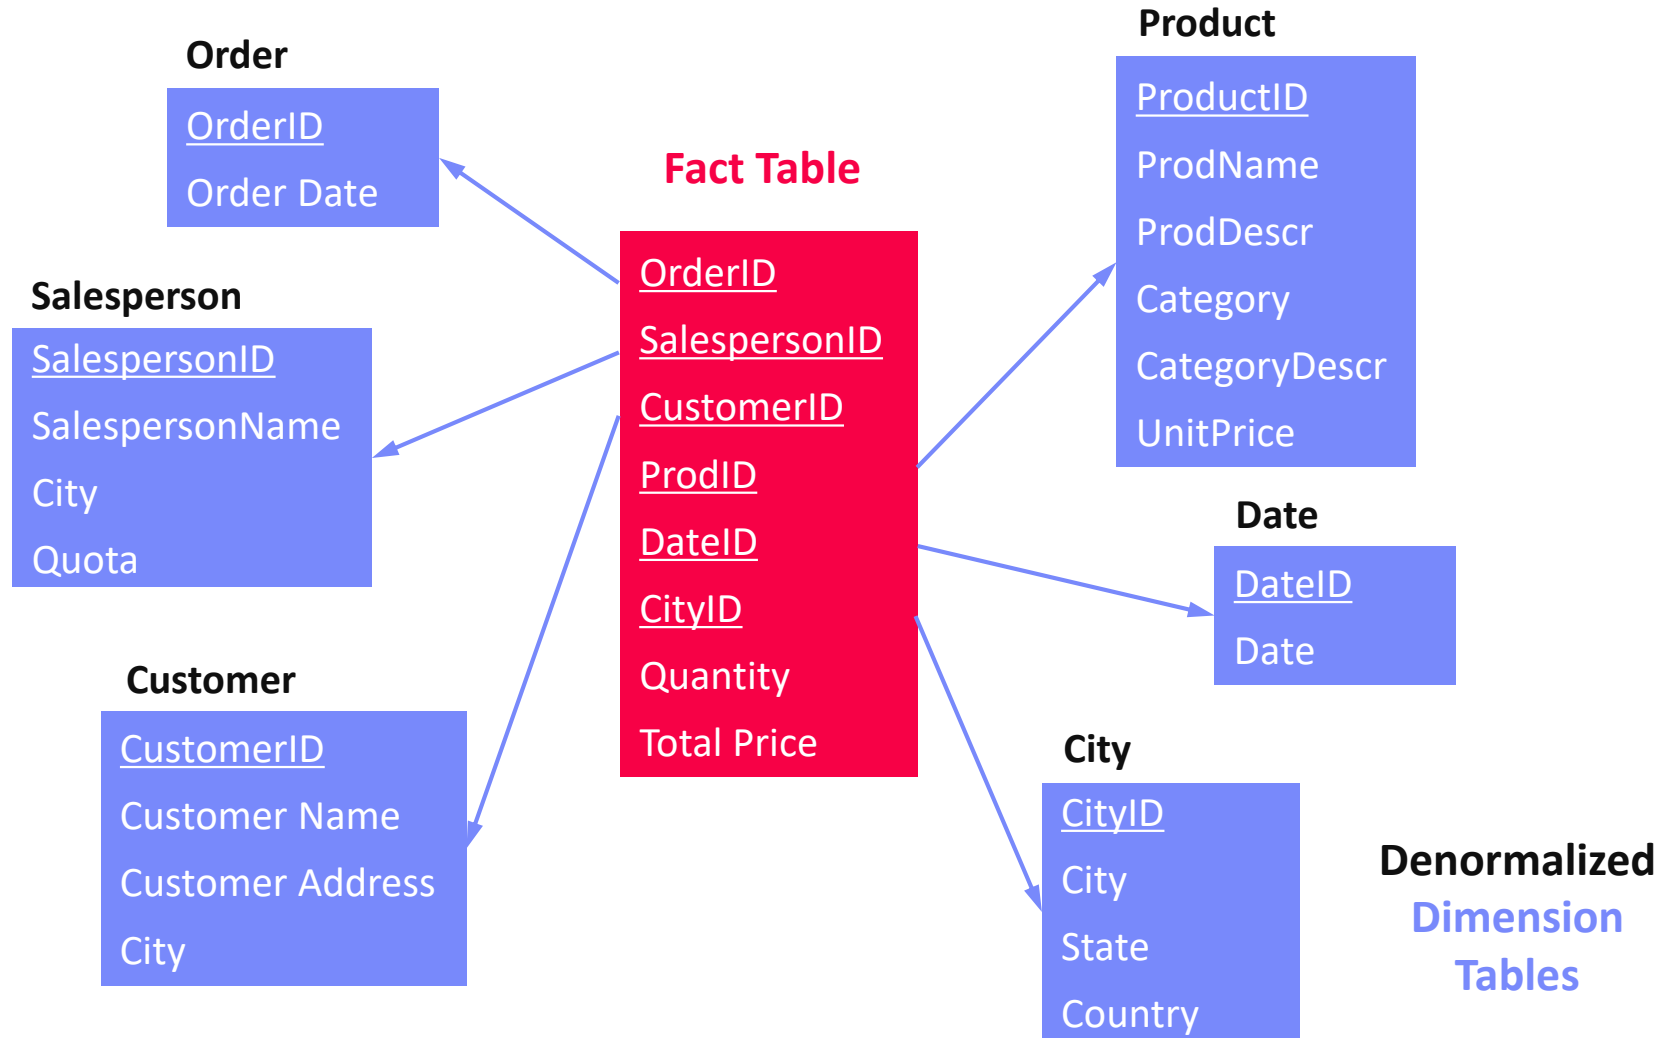
	A1 INT	A2 INT	A3 BOOL
	3	7	T
	1	2	T
	3	4	F
Tuple	1	7	T

cardinality: 4  
rank: 3

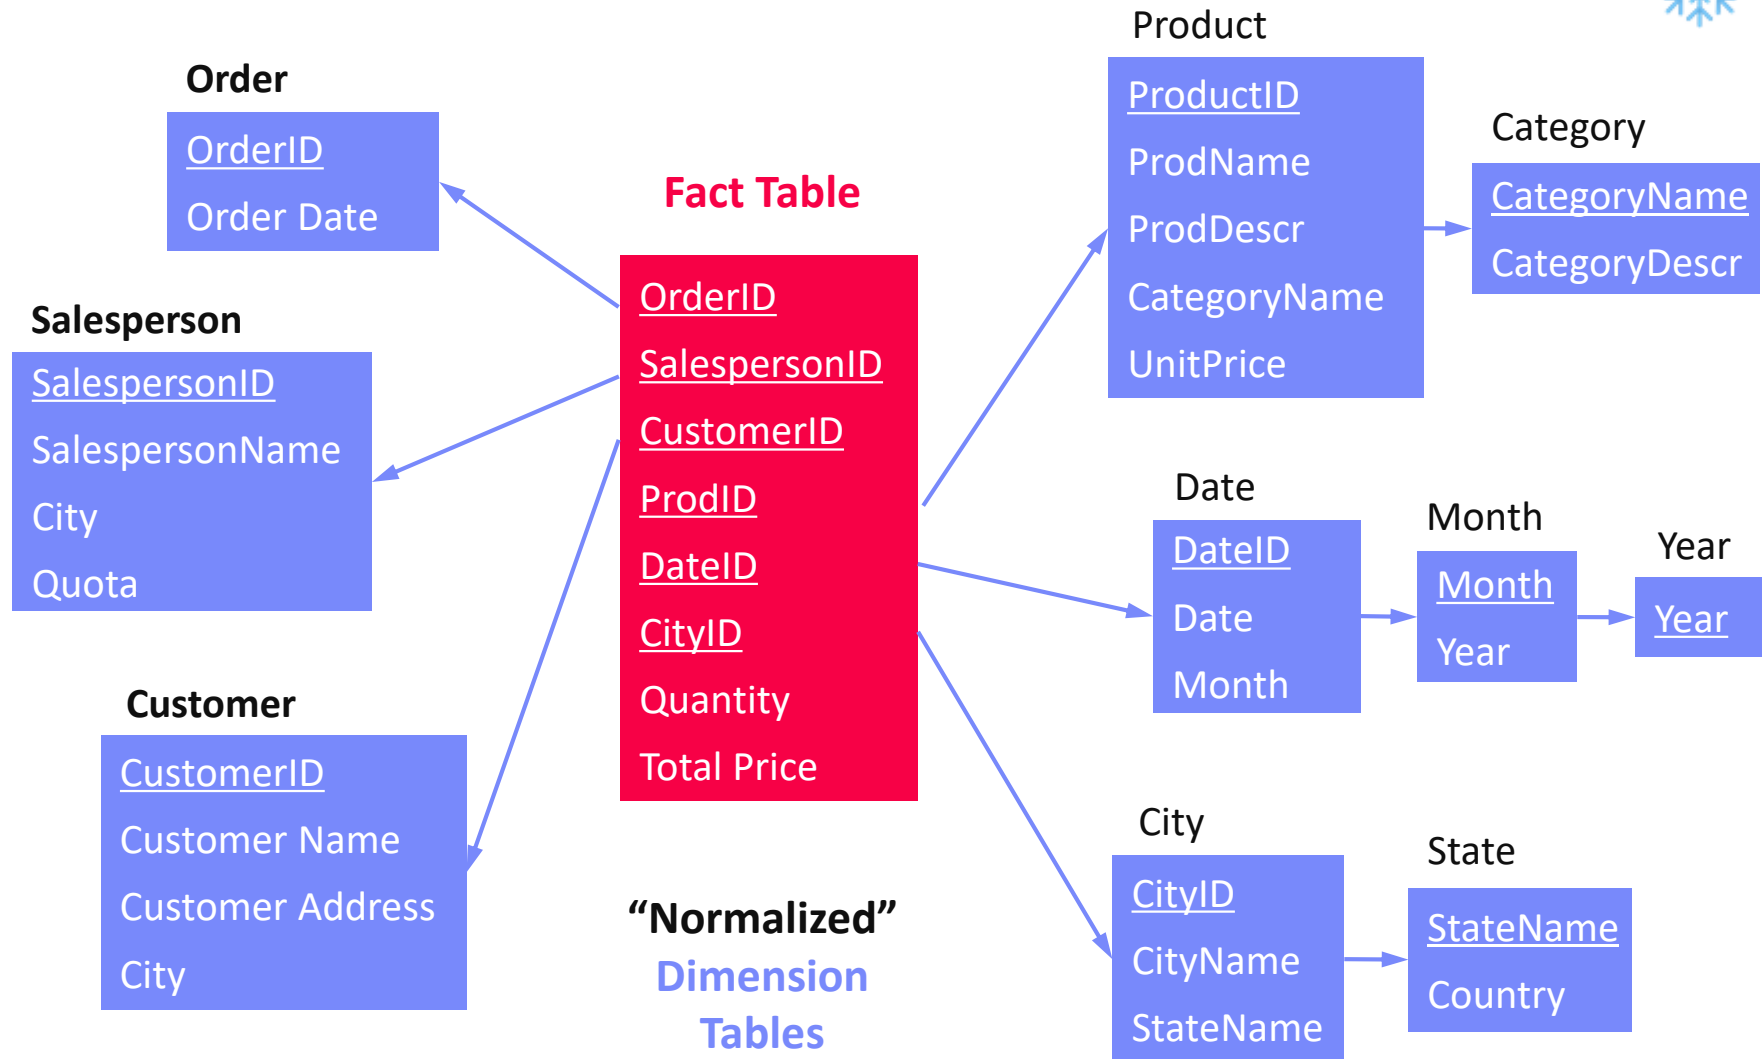
- Additional Terminology

- Tuple: row of k elements of a relation
- Cardinality of a relation: number of tuples in the relation
- Rank of a relation: number of attributes
- Semantics: Set := no duplicate tuples (in practice: Bag := duplicates allowed)
- Order of tuples and attributes is irrelevant

# ROLAP – Star Schema




# ROLAP – Snowflake Schema



# ROLAP – Other Schemas

## ■ Galaxy Schema

- Similar to **star**-schema but with **multiple fact tables** and potentially shared dimension tables
- Multiple stars → Galaxy



## ■ Snow-Storm Schema

- Similar to **snow-flake**-schema but with **multiple fact tables** and potentially shared dimension tables
- Multiple snow flakes → snow storm



## ■ OLAP Benchmark Schemas

- TPC-H** (8 tables, normalized schema)
- SSB** (5 tables, star schema, simplified TPC-H)
- TPC-DS** (24 tables, snow-storm schema)

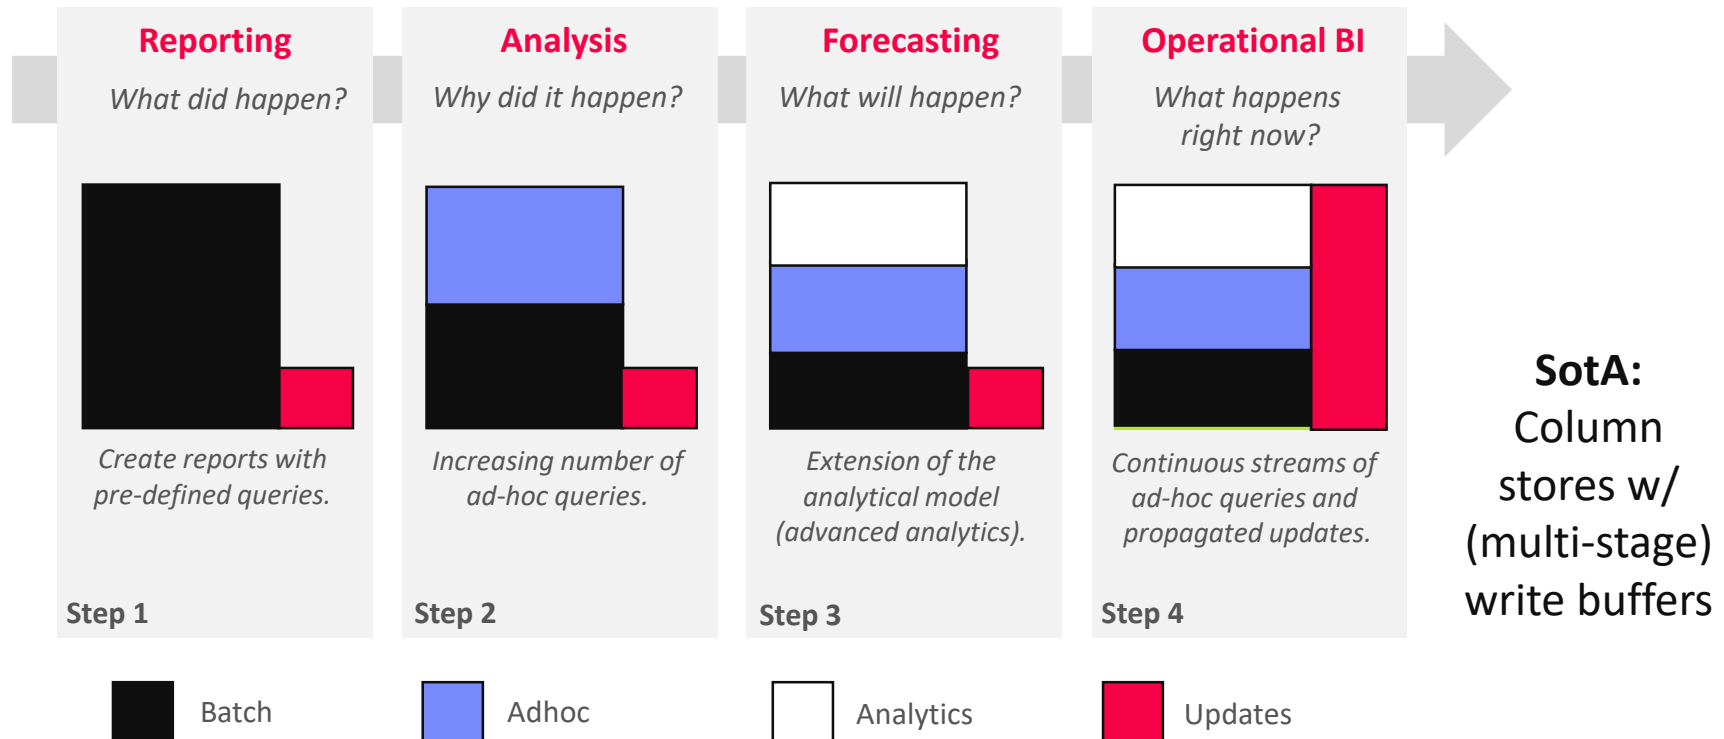
*“TPC-D and its successors, TPC-H and TPC-R assumed a 3rd Normal Form (3NF) schema. However, over the years the industry has expanded towards star schema approaches.”*

[Raghunath Othayoth Nambiar, Meikel Poess: The Making of TPC- DS. **VLDB 2006**]



# Evolution of DWH/OLAP Workloads

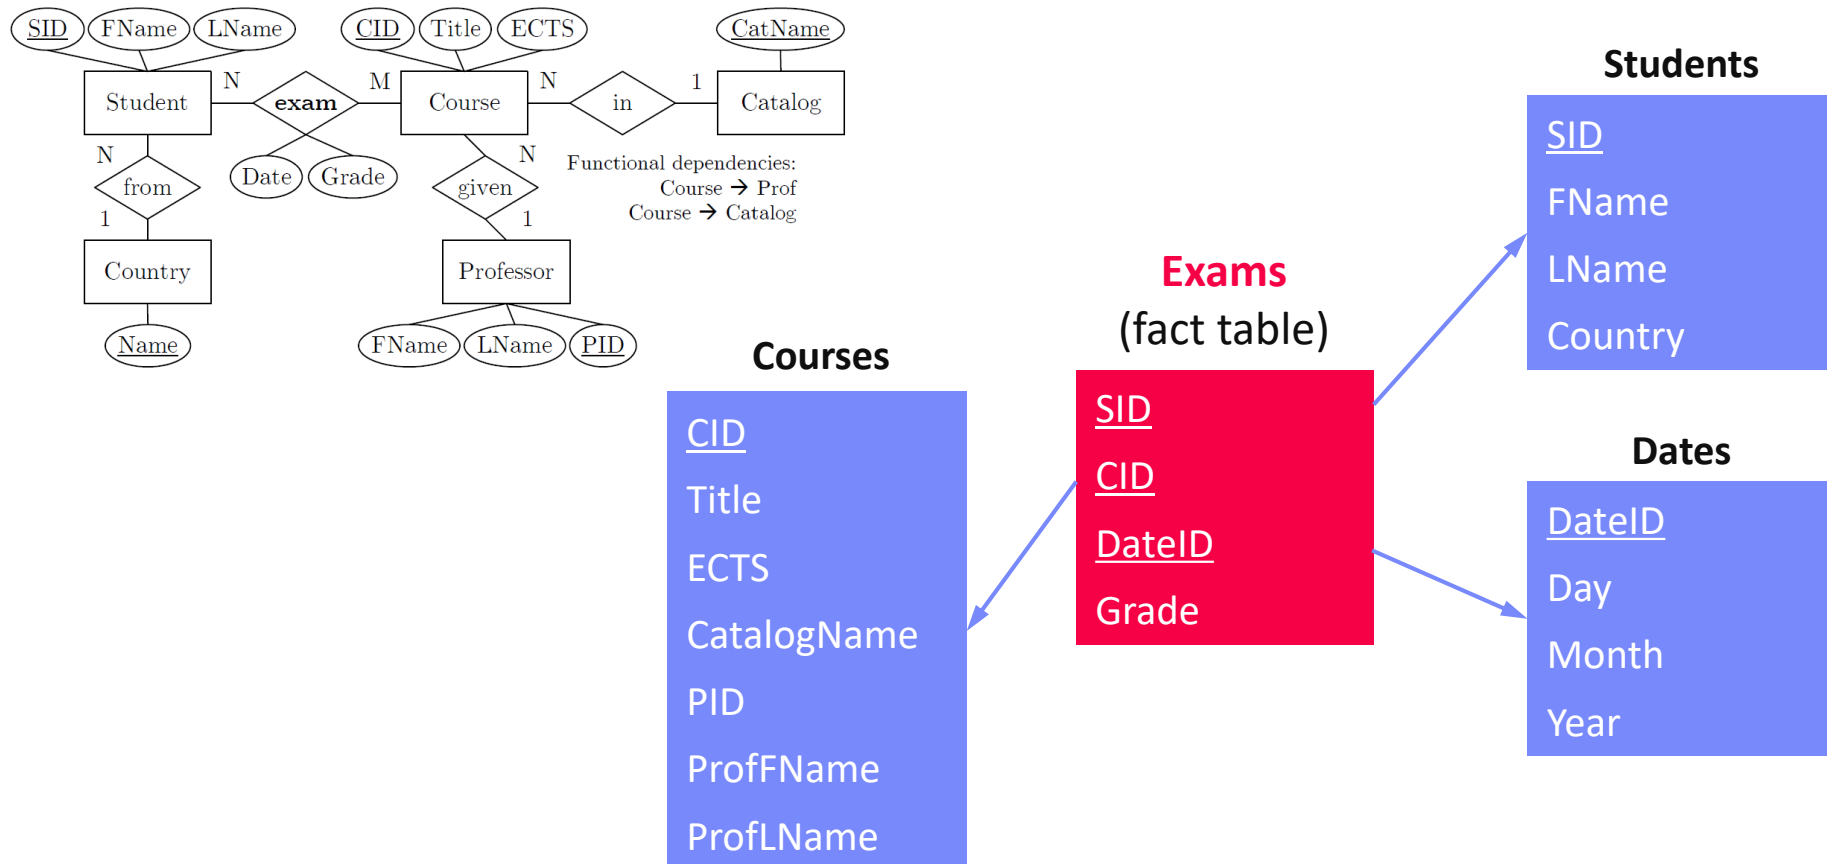
## ■ Goals: Advanced analytics and Operational BI



# BREAK (and Test Yourself)

[Exam Feb 08, 2021]

- Task: Given below ER diagram, create a ROLAP star schema. Data types can be ignored, but indicate PK and FK constraints. (9/100 points)



# Extraction, Transformation, Loading (ETL)

# Extract-Transform-Load (ETL) Overview

## ■ Overview

- ETL process refers to the overall process of obtaining data from the source systems, cleaning and transforming it, and loading it into the DWH
- Subsumes many integration and cleaning techniques

## ■ #1 ETL

- Extract data from heterogeneous sources
- Transform data via dedicated data flows or in staging area
- Load cleaned and transformed data into DWH

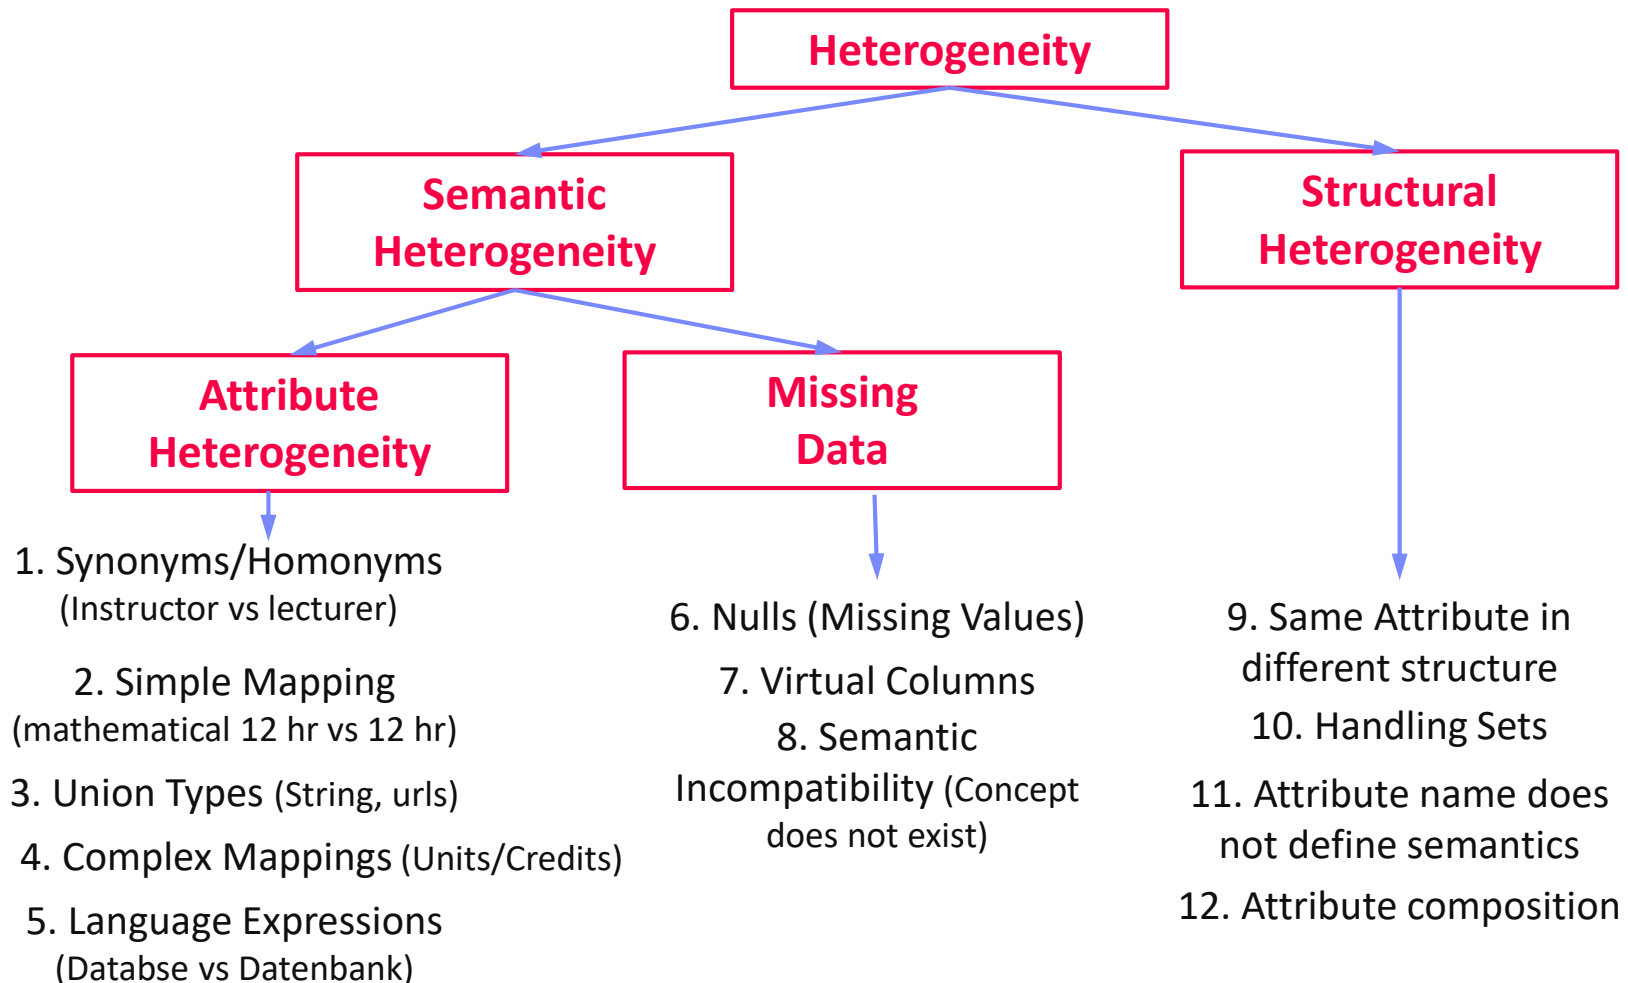
## ■ #2 ELT

- Extract data from heterogeneous sources
  - Load raw data directly into DWH
  - Perform data transformations inside the DWH via SQL
- ➔ allows for **automatic optimization of execution plans**



# Types of Heterogeneity

[J. Hammer, M. Stonebraker, and O. Topsakal:  
THALIA: Test Harness for the Assessment of  
Legacy Information Integration Approaches. U  
Florida, TR05-001, 2005]



# Corrupted Data

## ■ Heterogeneity of Data Sources

- Update anomalies on denormalized data / eventual consistency
- Changes of app/preprocessing over time (US vs us) → inconsistencies

## ■ Human Error

- Errors in semi-manual data collection, laziness (see default values), bias
- Errors in data labeling (especially if large-scale: crowd workers / users)

## ■ Measurement/Processing Errors

- Unreliable HW/SW and measurement equipment (e.g., batteries)
- Harsh environments (temperature, movement) → aging

**Uniqueness & duplicates**

**Contradictions & wrong values**

**Missing Values**

**Ref. Integrity**

[Credit: Felix Naumann]

ID	Name	BDay	Age	Sex	Phone	Zip
3	Smith, Jane	05/06/1975	44	F	999-9999	98120
3	John Smith	38/12/1963	55	M	867-4511	11111
7	Jane Smith	05/06/1975	24	F	567-3211	98120

Zip	City
98120	San Jose
90001	Lost Angeles

**Typos**

# ETL – Planning and Design Phase

## ■ Architecture, Flows, and Schemas

- #1 Plan requirements, architecture, tools
- #2 Design high-level integration flows (systems, integration jobs)
- #3 Data understanding (copy/code books, meta data)
- #4 Design dimension loading (static, dynamic incl keys)
- #5 Design fact table loading

## ■ Data Integration and Cleaning

- #5 Types of data sources (snapshot, APIs, query language, logs)
- #6 Prepare schema mappings → see [04 Schema Matching and Mapping](#)
- #7 Change data capture and incremental loading (diff, aggregates)
- #8 Transformations, enrichments, and deduplication → [05 Entity Linking](#)
- #9 Data validation and cleansing → see [06 Data Cleaning and Data Fusion](#)

## ■ Optimization


- #10 Partitioning schemes for loaded data (e.g., per month)
- #11 Materialized views and incremental maintenance

# Events and Change Data Capture


- **Goal: Monitoring operations of data sources for detecting changes**
- **#1 Explicit Messages/Triggers**
  - Setup update propagation from the source systems to middleware
  - Asynchronously propagate the updates into the DWH
- **#2 Log-based Capture**
  - Parse system logs / provenance to retrieve changes since last loading
  - Sometimes combined w/ replication → **03 MoM, EAI, and Replication**
  - Leverage explicit audit columns or internal timestamps
- **#3 Snapshot Differences**
  - Compute difference between old and new snapshot (e.g., files) before loading
  - Broadly applicable but more expensive

# Example ETL Flow

- **Example Flows**  
(**Pentaho Data Integration**, since 2015 Hitachi)



[Matthias Boehm, Uwe Wloka, Dirk Habich, Wolfgang Lehner: GCIP: exploiting the generation and optimization of integration processes. **EDBT 2009**]



[Alkis Simitsis, Kevin Wilkinson, Petar Jovanovic: xPAD: a platform for analytic data flows. **SIGMOD 2013**]



- **Other Tools**

- IBM IS, Informatica, SAP BO, MS Integration Services
- Open Source: Pentaho Data Integration, Scriptella ETL, CloverETL, Talend

# ETL via Apache Spark

## ■ Example

- Distributed ETL pipeline processing

[Xiao Li: Building Robust ETL  
Pipelines with Apache Spark,  
**Spark Summit 2017**]



**//load csv and postgres tables**

```
val csvTable = spark.read.csv("/source/path")
val jdbcTable = spark.read.format("jdbc")
    .option("url", "jdbc:postgresql:...")
    .option("dbtable", "TEST.PEOPLE")
    .load()
```

**//join tables, filter and write as parquet**

```
csvTable
    .join(jdbcTable, Seq("name"), "outer")
    .filter("id <= 2999")
    .write.mode("overwrite")
    .format("parquet")
    .saveAsTable("outputTableName")
```

**11 Distributed, Data-  
Parallel Computation**

# SQL/OLAP Extensions

# Overview Multi-Groupings

## Recap: GROUP BY


- Group tuples by categorical variables
- Aggregate per group

Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

```
SELECT Year, SUM(Revenue)
FROM Sales
GROUP BY Year
```

Year	SUM
2004	60
2005	30

## Grouping Extensions





# Grouping Sets

**GROUP BY GROUPING SETS**  
(((<attribute-list>)), ...)

## Semantics

- Grouping by multiple group-by attribute lists w/ consistent agg function
- Equivalent to multiple GROUP BY, connected by UNION ALL

## Example

```
SELECT Year, Quarter, SUM(Revenue)
FROM R
GROUP BY GROUPING SETS
((), (Year), (Year,Quarter))
```

Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30



Year	Quarter	SUM
-	-	90
2004	-	60
2005	-	30
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

# Rollup (see also multi-dim ops)

**GROUP BY ROLLUP**  
(**<attribute-list>**)

## Semantics

- Hierarchical grouping along dimension hierarchy
- GROUP BY ROLLUP (A1,A2,A3)**  
:= GROUP BY GROUPING SETS((), (A1), (A1,A2), (A1,A2,A3))

## Example

```
SELECT Year, Quarter, SUM(Revenue)
FROM R
GROUP BY ROLLUP(Year,Quarter)
```

Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30




Year	Quarter	SUM
-	-	90
2004	-	60
2005	-	30
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

# Rollup, cont. and Grouping

## Operator Implementation

- Aggregation towers for (semi-)additive aggregation functions
- Example

```
SELECT Year, Quarter, SUM(Revenue)
FROM R
GROUP BY ROLLUP(Year, Quarter)
```



## GROUPING Semantics

- With ROLLUP or CUBE to identify aggregates
- NULL group vs NULL due to aggregation
- Example

```
SELECT Team, SUM(Revenue),
       GROUPING(Team) AS Agg
FROM R
GROUP BY ROLLUP (Team)
```

Team	Revenue	Agg
NULL	10	0
Sales	40	0
Tech	20	0
NULL	70	1

# Cube

GROUP BY **CUBE**(<attribute-list>)

## Semantics

- Computes aggregate for all  $2^n$  combinations for  $n$  grouping attributes
- Equivalent to enumeration via GROUPING SETS

## Example

```
SELECT Year, Quarter, SUM(Revenue)
FROM R
GROUP BY CUBE(Year,Quarter)
```

Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30



Year	Quarter	SUM
-	-	90
2004	-	60
2005	-	30
-	1	40
-	2	20
-	3	10
-	4	20
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

# Cube, cont.

## ■ Operator Implementation


- **Aggregation lattice** for (semi-)additive aggregation functions
- **But: multiple alternative paths**  
→ how to select the cheapest?

## ■ Recap: Physical Group-By Operators

- SortGroupBy / -Aggregate
- HashGroupBy / -Aggregate

## ■ Cube Implementation Strategies

- #1: Some operators can share sorted order (e.g.,  $\{A,B\} \rightarrow \{A\}$ )
- #2: Subsets with different cardinality → pick smallest intermediates



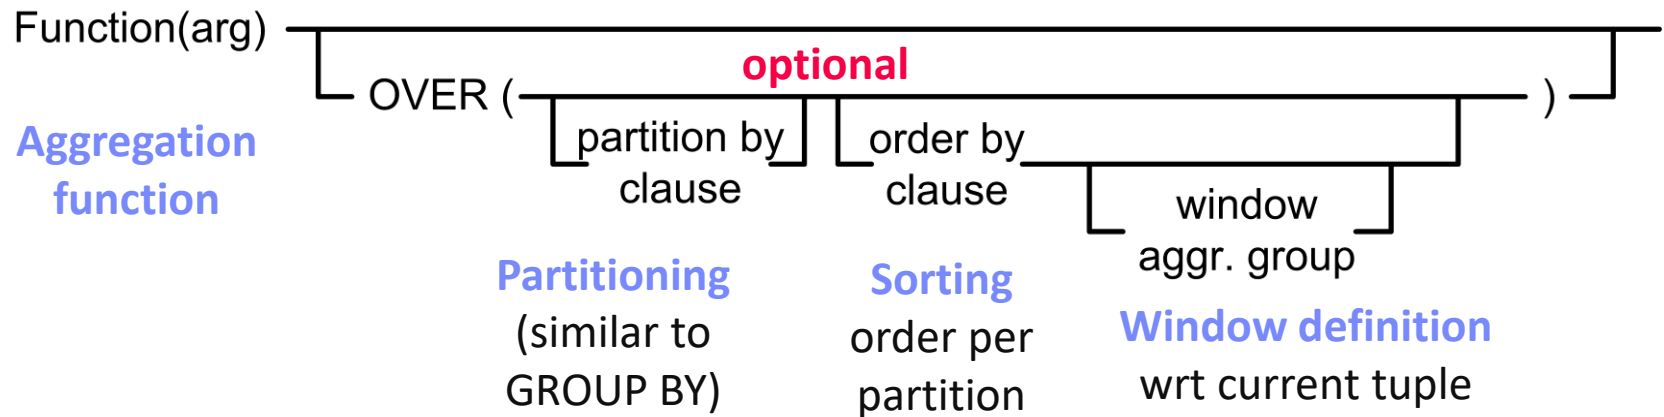
# Overview Reporting Functions

## ■ Motivation and Problem


- Scalar functions as well as grouping + aggregation
- For many advanced use cases **not flexible enough**

## ■ Reporting Functions

- Separate partitioning (grouping) and aggregation via OVER
- Allows local partitioning via windows and ranking/numbering



# RF – Aggregation Function



## Semantics

- Operates over window and returns value for every tuple
- RANK(), DENSE\_RANK(), PERCENT\_RANK(), CUME\_DIST(), ROW\_NUMBER()

## Example

```

SELECT Year, Quarter,
       RANK() OVER (ORDER BY Revenue ASC) AS Rank1,
       DENSE_RANK() OVER (ORDER BY Revenue ASC) AS DRank1,
FROM R

```


Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30



OVER()  
represents  
all tuples

Year	Quarter	Rank1	DRank1
2004	1	1	1
2004	3	1	1
2004	2	3	2
2004	4	3	2
2005	1	5	3

# RF – Partitioning



## ■ Semantics

- Select tuples for aggregation via **PARTITION BY** <attribute-list>

## ■ Example

```
SELECT Year, Quarter, Revenue,
       SUM(Revenue) OVER(PARTITION BY Year)
FROM R
```


Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30



Year	Quarter	Revenue	SUM
2004	1	10	60
2004	2	20	60
2004	3	10	60
2004	4	20	60
2005	1	30	30



# RF – Partition Sorting



## Semantics

- Define computation per partition via **ORDER BY** <attribute-list>
- Note: ORDER BY allows cumulative computation → cumsum()



## Example

```
SELECT Year, Quarter, Revenue,
       SUM(Revenue) OVER(PARTITION BY Year ORDER BY Quarter)
FROM R
```

Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30



Year	Quarter	Revenue	SUM
2004	1	10	10
2004	2	20	30
2004	3	10	40
2004	4	20	60
2005	1	30	30

# 42 RF – Windowing

## ■ Semantics

- Define window for computation (e.g., for moving average, cumsum)


## ■ Example

```
SELECT Year, Quarter, Revenue, AVG(Revenue)
OVER (ORDER BY Year, Quarter
ROWS BETWEEN 1 PRECEDING AND CURRENT ROW)
FROM R
```


Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30



Year	Quarter	Revenue	AVG
2004	1	10	10
2004	2	20	15
2004	3	10	15
2004	4	20	15
2005	1	30	25



## Measurements



# Trend: Cloud Data Warehousing

## 10 Distributed Data Storage

- #1 **Google** Big Query

[Google, Kazunori Sato: An Inside Look at Google BigQuery, Google **White Paper** 2012]



- #2 **Amazon** Redshift

[Anurag Gupta, Deepak Agarwal, Derek Tan, Jakub Kulesza, Rahul Pathak, Stefano Stefani, Vidhya Srinivasan: Amazon Redshift and the Case for Simpler **Data Warehouses**. **SIGMOD** 2015]



- #3 **Microsoft** Azure Data Warehouse

- #4 **IBM** BlueMix dashDB

[IBM: IBM dashDB - Cloud-based **data warehousing** as-a-service, built for analytics, IBM **White Paper** 2015]



- #5 **Snowflake** Data Warehouse

[Benoît Dageville et al.: The Snowflake Elastic **Data Warehouse**. **SIGMOD** 2016]




# Trend: Data Lakes and Lakehouse

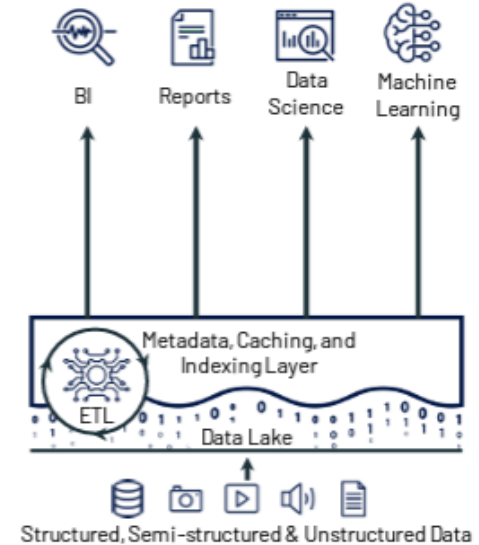
10 Distributed  
Data Storage



(a) First-generation platforms.



(b) Current two-tier architectures.



(c) Lakehouse platforms.

[Matei Zahari et. al, Lakehouse: A New Generation of Open Platforms that Unify Data Warehousing and Advanced Analytics. **CIDR 2021**]



[Alkis Simitsis et. al., The History, Present, and Future of ETL Technology, **DOLAP 2023**]



# Summary and Q&A

## ■ Data Warehousing (DWH)

- DWH architecture
- Multidimensional modeling

## ■ Extraction, Transformation, Loading (ETL)


- ETL process, errors, and data flows

## ■ SQL/OLAP Extensions

- Multi-grouping operations
- Reporting functions

## ■ Next Lectures (**Data Integration Architectures**)

- **03 Message-oriented Middleware, EAI, and Replication** [Oct 21]
- **04 Schema Matching and Mapping** [Oct 28]
- **05 Entity Linking and Deduplication** [Nov 04]
- **06 Data Cleaning and Data Fusion** [Nov 11]



“There is a profound cultural assumption in the business world that *if only we could see all of our data, we could manage our businesses more effectively*. This cultural assumption is so deeply rooted that we take it for granted. Yet this is the mission of the data warehouse, and this is why *the data warehouse is a permanent entity [...] even as it morphs and changes its shape.*”  
-- Ralph Kimball, Joe Caserta;  
**2004**